

FirmScope

An AI agent that audits a law firm's website and writes the email you would send them.

WHAT IT IS

You paste in a US law firm's website. About two minutes later you get a score out of 100, a list of what is wrong with the site (with proof for each one), and a cold email you could actually send to the managing partner.

WHY I BUILT IT

I do SEO for law firms. Before you can pitch one, somebody has to go through their whole site by hand, work out what is actually losing them cases, and write an email specific enough that a partner bothers replying. That takes about an hour per firm, and it is the thing that slows down outreach more than anything else. This does the same pass in two minutes.

HOW YOU USE IT

1. Paste a law firm URL and hit run.
2. Watch it work. Every step it takes shows up live, so you can see it planning, crawling, checking, and writing.
3. Read the report, copy the email, done. It all gets saved so you can come back to it.

The one idea behind the whole thing: tools measure, the model judges.

Every number in the report comes from code, never from the AI. Word counts, missing tags, page speed, how many different phone numbers are on the site. The AI is not allowed to state a fact that a tool did not hand it. What the AI does instead is the part code cannot do: read the actual writing on the page and decide whether it is any good, then explain what it costs the firm and write it up. That is what keeps the report trustworthy enough to put in front of a client.

WHAT IS HAPPENING UNDER THE HOOD

- It crawls the site and picks a sample of the pages that matter (practice areas, lawyer bios, contact, results).
- Around 32 checks run in plain code and produce pass or fail results.
- A helper agent reads the real page copy and judges whether it is written for that state and that firm, or is generic filler.
- A scoring function turns all of it into a number. The AI cannot move that number up or down.
- A second helper agent writes the email, leading with the single worst thing it found.

BUILT WITH

Next.js, Tailwind, shadcn/ui, Supabase for the database, LangGraph deep agents for the harness, Gemini for the model, deployed on Vercel.

WHAT I WOULD ADD NEXT

Compare the firm against the ones ranking above them, PDF export of the report, and a batch mode so you can feed it a whole list of firms at once.

THINGS THAT BROKE, AND GOT FIXED

- It told a firm with 47 lawyer bios that it had none. It was only looking at the handful of pages it downloaded instead of the whole site.
- The email made up a competitor stat. It said "competitors ranking in your market average 800 to 1,500 words" when the tool has no competitor data at all. That is the worst kind of bug for something meant to be trustworthy, so it is now blocked outright.
- It crashed on the live site while working perfectly on my machine. Turned out the agent library was not getting packaged into the deployment.